



McGill

ECSE 421 Lecture 2: Specifying Requirements Early Design Stages

ESD Chapter 2

© Peter Marwedel, Brett H. Meyer

Last Time

- Introduction to Embedded System Design
 - Embedded System
 - Cyber-physical system
 - Real-time system
- Common characteristics
- Problem: *design challenges*
- Solution: *design flow*

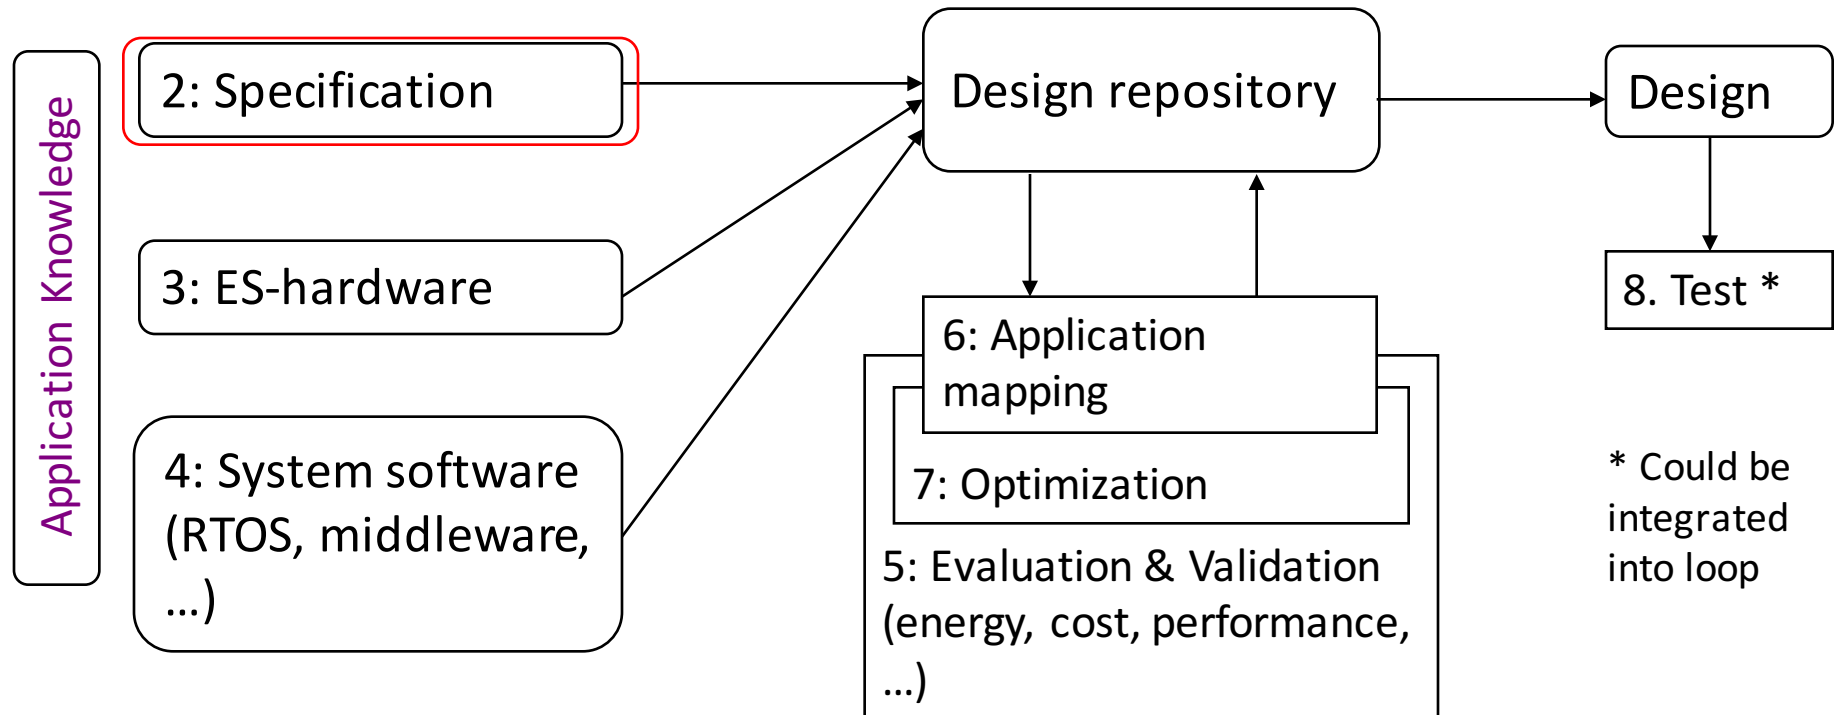
Where Are We?

W	D	Date		Topic	ESD	Out	In	Notes
1	T	12-Jan-2016	L01	Introduction to Embedded System Design	1.1-1.4			
	R	14-Jan-2016	L02	Specifying Requirements / MoCs / MSC	2.1-2.3			
2	T	19-Jan-2016	L03	CFSMs	2.4			
	R	21-Jan-2016	L04	Data Flow Modeling	2.5			
3	T	26-Jan-2016	L05	Petri Nets	2.6			
	R	28-Jan-2016	L06	Discrete Event Models	2.7			
4	T	2-Feb-2016	L07	DES / Von Neumann Model of Computation	2.8-2.10			Guest: Z. Al-bayati
	R	4-Feb-2016	L08	Sensors	3.1-3.2			
5	T	9-Feb-2016	L09	Processing Elements	3.3			
	R	11-Feb-2016		Processing Elements	3.3			
6	T	16-Feb-2016	L10	More Processing Elements / FPGAs				
	R	18-Feb-2016	L11	Memories, Communication, Output	3.4-3.6			
7	T	23-Feb-2016	L12	Embedded Operating Systems	4.1			
	R	25-Feb-2016		Midterm exam: in-class, closed book				Chapters 1-3
8	T	2-Mar-2016		No class				Winter break
	R	4-Mar-2016		No class				Winter break
	T	9-Mar-2016	L13	Middleware	4.4-4.5			
	R	11-Mar-2016	L14	Performance Evaluation	5.1-5.2			
9	T	16-Mar-2016	L15	More Evaluation and Validation	5.3-5.8			
	R	18-Mar-2016	L16	Introduction to Scheduling	6.1-6.2.2			
10	T	23-Mar-2016	L17	Scheduling Aperiodic Tasks	6.2.3-6.2.4			
	R	25-Mar-2016	L18	Scheduling Periodic Tasks	6.2.5-6.2.6			

Today

- Requirements \Rightarrow Specification
 - What do we need to specify?
 - How can we specify it?
- Models of Computation
 - What does it mean to “compute”?
 - Is there more than one way to describe “computation”?
- Modeling in early design stages

Hypothetical Design Flow



Numbers denote sequence of chapters

System Specifications

- Why bother with specification?
- If there's an error in the specification
 - It will be difficult to get the design right
 - This potentially wastes a lot of time
- Typically, we work with **models** of the **system under design** (SUD)
- What is a model anyway?




Models

Definition: *A model is a simplification of another entity ... The model contains exactly those characteristics and properties of the modeled entity that are relevant for a given task. A model is minimal with respect to a task if it does not contain any other characteristics than those relevant for the task. [Jantsch, 2004]*


What requirements do we have for our models?

Spec. Req's: Hierarchy

- Humans incapable of understanding systems with more than ~5 objects
- Most actual systems require more objects
⇒ Hierarchy
 - Behavioral hierarchy
Examples: states, processes, procedures
 - Structural hierarchy
Examples: processors, racks, printed circuit boards




proc
proc
proc



Spec. Req's: Component-based design

- Systems must be designed from components
- Must be “easy” to derive system behavior from the behavior of subsystems
- Must provide support for
 - Concurrency
 - Synchronization and communication
- Work of Sifakis, Thiele, Ernst, ...



Spec. Req's: Timing

- Modeling application timing behavior is essential for embedded and cy-phy systems!
 - Additional information (periods, dependences, scenarios, use cases) is helpful, too!
- Furthermore, the performance of the underlying platform must be known
 - And often must be predictable
- This has far-reaching consequences for design processes!




“The lack of timing in the core abstraction (of computer science) is a flaw, from the perspective of embedded software” [Lee, 2005]

Specifying Timing


- Four types of techniques are required for specifying timing requirements [Burns, 1990]

1. Measure elapsed time

E.g., to check, how much time has elapsed since last call




2. Delay processes





Specifying Timing, Cont'd

3. Specify timeouts

E.g., to stay in a certain state a maximum time




4. Specifying deadlines and schedules



Spec. Req's: Supporting Reactive Systems

- *State-oriented behavior*
 - Required for reactive systems; classical automata are insufficient
- *Event-handling*
 - External or internal events
- *Exception-oriented behavior*
 - Describing exceptions for every state is inefficient




We will see how all the arrows labeled *k* can be replaced by a single one.

Additional Specification Requirements

- Presence of programming elements
- Executability (no algebraic specification)
- Support for the design of large systems (\Rightarrow OO)
- Domain-specific support
- Readability
- Portability and flexibility
- Termination
- Support for non-standard I/O devices
- Non-functional properties
- Support for the design of dependable systems
- No obstacles for efficient implementation
- Adequate model of computation

Models of Computation

- **What does it mean “to compute”?**
- Models of computation define:
 - Components and an execution model for computations for each component
 - Communication model for exchange of information between components



Why Not von Neumann?



- Data races occur when accessing shared resources

```
task a {
```

```
..
```

```
P(S) //obtain lock
```

```
.. // critical section
```

```
V(S) //release lock
```

```
}
```

```
task b {
```

```
..
```

```
P(S) //obtain lock
```

```
.. // critical section
```

```
V(S) //release lock
```

```
}
```

Race-free access to shared memory protected by S is possible

- In general-purpose computing, we address this
 - With synchronization primitives (semaphores)
 - With cache coherence protocols


Challenges With the von Neumann MoC

- Thread-based multiprocessing may access global variables
- Operating systems theory tells us that
 - access to global variables may lead to race conditions,
 - to avoid these, we need to use mutual exclusion,
 - mutual exclusion may lead to deadlocks, and
 - avoiding deadlocks is possible only if we accept performance penalties

Avoiding deadlock is hard in practice; more on this later

Dependence Graphs: Definition

Sequence
constraint




Nodes could be programs or simple operations

- **Def.:** A **dependence graph** is a directed graph $G=(V,E)$ in which $E \subseteq V \times V$ is a relation.
- If $(v_1, v_2) \in E$, then v_1 is called an **immediate predecessor** of v_2 and v_2 is called an **immediate successor** of v_1 .
- Suppose E^* is the transitive closure of E . If $(v_1, v_2) \in E^*$, then v_1 is called a **predecessor** of v_2 and v_2 is called a **successor** of v_1 .


Dependence Graphs: Timing Information

- Dependence graphs may contain additional information, for example, timing information





Dependence Graphs: I/O-information

- Nodes can also be used to capture information about inputs and outputs




Dependence Graphs: Shared Resources

- Additional information can capture which tasks use shared resources (i.e., structural hazards)




Dependence Graphs: Periodic Schedules

- A job is a single execution of the dependence graph
- Periodic dependence graphs are infinite



Dependence Graphs: Hierarchy

- Dependence graphs can be hierarchical
- This makes it easier to reason about complex graphs both for
 - Human developers, and
 - Algorithms



Modeling Communication

- Two basic approaches
 - Shared memory
 - Message passing
- Shared memory
 - Variables are accessible to several components/tasks



Shared Memory



- Data races are possible
 - Leading to potentially inconsistent results
- Exclusive access to critical sections must be guaranteed

```
task a {  
  ..  
  P(S) //obtain lock  
  .. // critical section  
  V(S) //release lock  
}
```


```
task b {  
  ..  
  P(S) //obtain lock  
  .. // critical section  
  V(S) //release lock  
}
```

Race-free access to shared memory protected by S is possible

P(S) and V(S) are **semaphore** operations,
allowing at most n accesses, $n = 1$ in this case (mutex, lock)

Non-blocking Message Passing


- Asynchronous message passing
- Sender does not have to wait until the message has arrived



- Potential problem: buffer overflow

Blocking Message Passing


- Synchronous message passing; rendez-vous
- Sender will wait until receiver has received the message



- No buffer overflow, but reduced performance

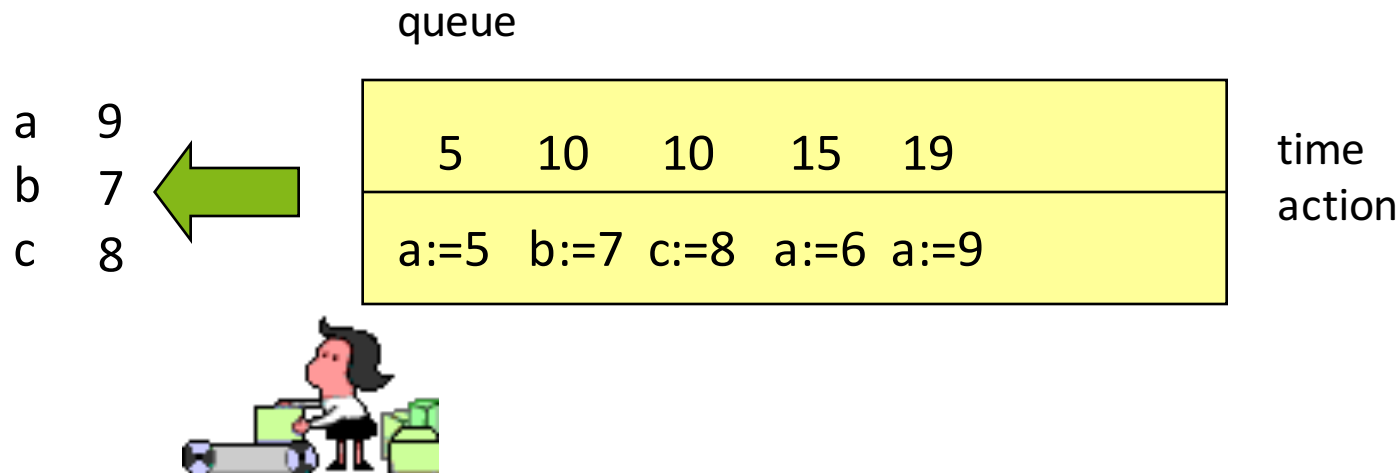
Organizing Computation

- Von Neumann
 - Sequential execution, program memory, etc.
- Finite state machines
 - Model the flow of execution
- Data flow
 - Models the flow of data
- Petri nets
 - Model synchronization



Organizing Computation, Cont'd

- Discrete event simulation



- Differential equations

$$\frac{\partial^2 x}{\partial t^2} = b$$



MoCs Considered in 421

Communication/ local computations	Shared memory	Message passing	
		Synchronous	Asynchronous
Undefined components		Plain text, use cases (Message) sequence charts	
Communicating finite state machines	StateCharts		SDL
Data flow	(Not useful)		Kahn networks, SDF
Petri nets		C/E nets, P/T nets, ...	
Discrete event (DE) model	VHDL*, Verilog*, SystemC*, ...	Only experimental systems, e.g. distributed DE in Ptolemy	
Von Neumann model	C, C++, Java	C, C++, Java with libraries CSP, ADA	

* Classification based on implementation with centralized data structures

Combined Models

- Specification and Description Language (SDL)
 - FSM+asynchronous message passing
- StateCharts
 - FSM+shared memory
- Communicating Sequential Processes (CSP); ADA
 - von Neumann execution+synchronous message passing
-
- See also
 - Work by Edward A. Lee, UCB
 - Axel Jantsch: *Modeling Embedded Systems and SoCs: Concurrency and Time in Models of Computation*, Morgan-Kaufman, 2004


Ptolemy

- Ptolemy (UC Berkeley) is an environment for simulating multiple models of computation
 - <http://ptolemy.berkeley.edu/>



Facing reality

- No language meets all requirements
⇒ Designers find the best compromise!



MoCs Considered in 421

Communication/ local computations	Shared memory	Message passing	
		Synchronous	Asynchronous
Undefined components		Plain text, use cases (Message) sequence charts	
Communicating finite state machines	StateCharts		SDL
Data flow	(Not useful)		Kahn networks, SDF
Petri nets		C/E nets, P/T nets, ...	
Discrete event (DE) model	VHDL*, Verilog*, SystemC*, ...	Only experimental systems, e.g. distributed DE in Ptolemy	
Von Neumann model	C, C++, Java	C, C++, Java with libraries CSP, ADA	

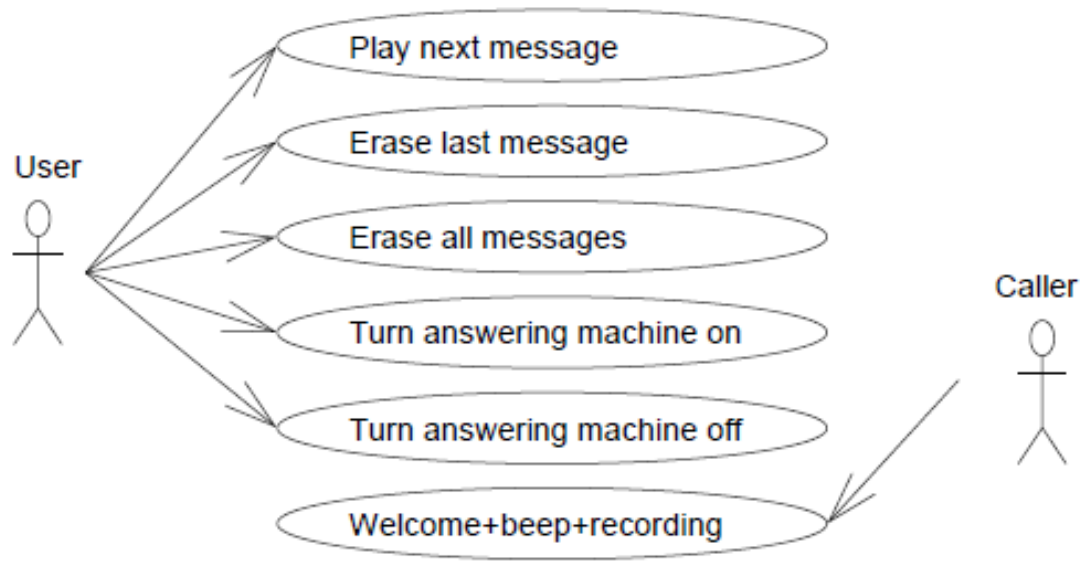
* Classification based on implementation with centralized data structures

Early Design Stages

- Capturing the requirements as text
 - Describe the system under design (SUD) in natural language
- Expectations for tools
 - Machine-readable
 - Version management
 - Dependency analysis
 - Example: DOORS® [Telelogic/IBM]
- How should requirements be captured?
 - Use cases
 - Message sequences

Use Cases

- Possible applications of the SUD
- Included in UML (Unified Modeling Language)
- Example: Answering machine

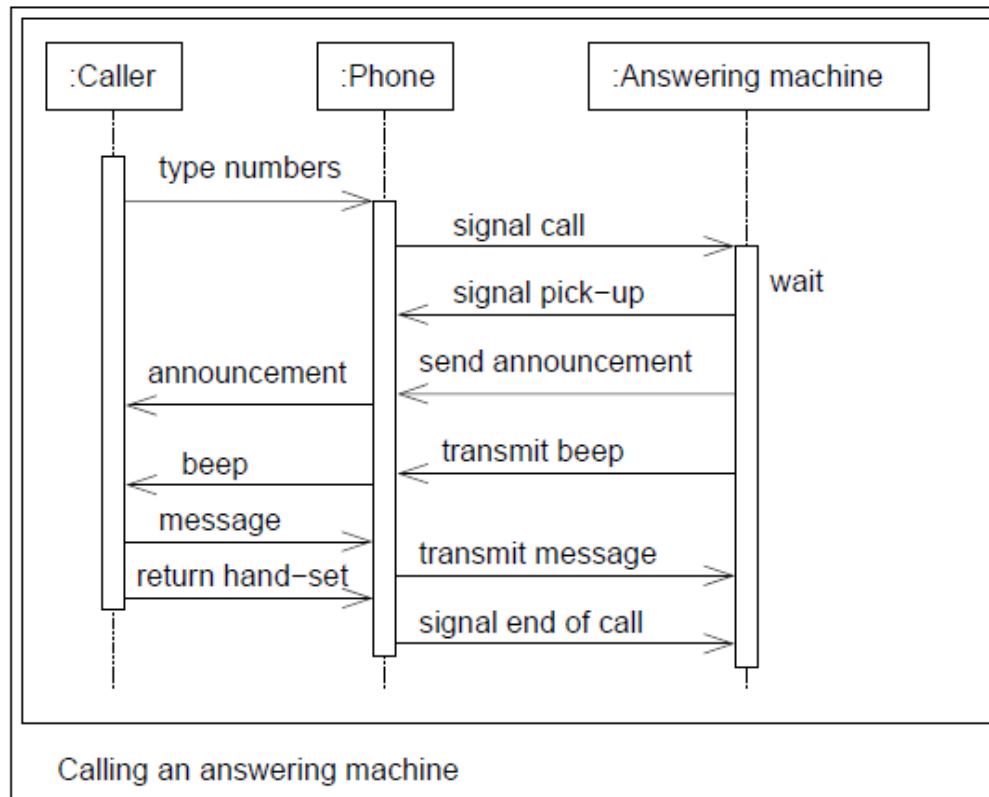


- Neither a precisely specified model of the computations nor a precisely specified model of the communication

(Message) Sequence Charts (MSC)

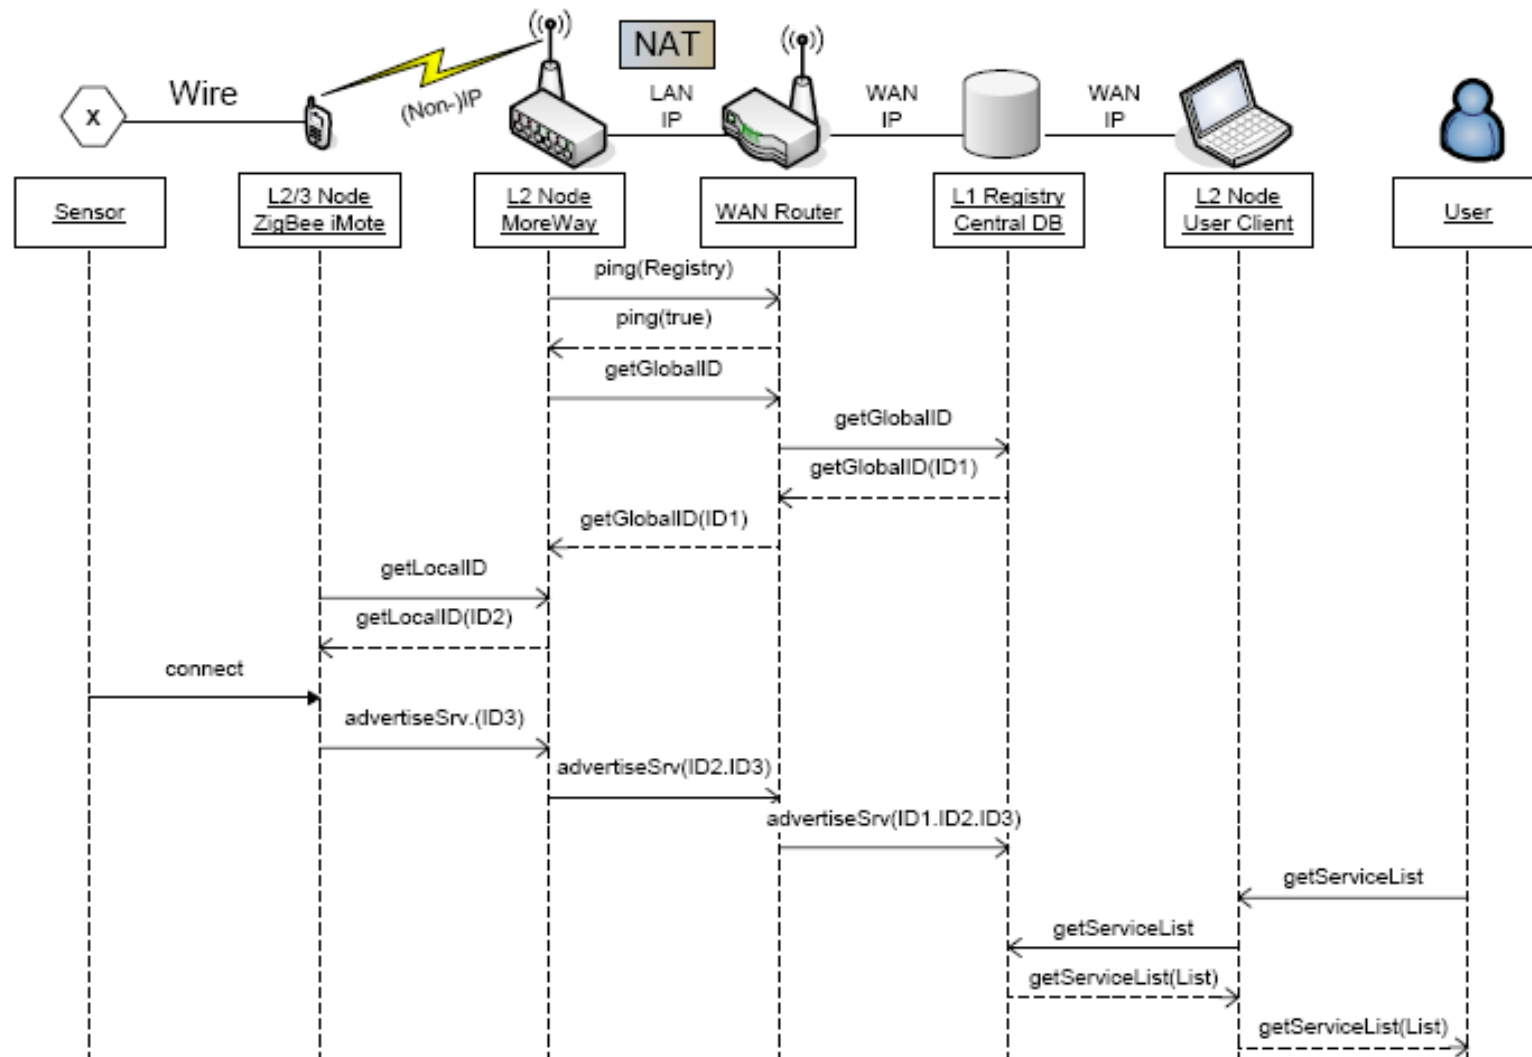
- Explicitly indicate exchange of information
- One dimension (usually vertical dimension) reflects time
- The other reflects distribution in space

Example:



- Included in UML
- Earlier called Message Sequence Charts, now mostly called Sequence Charts

Example: Sensor Services



www.ist-more.org, deliverable 2.1

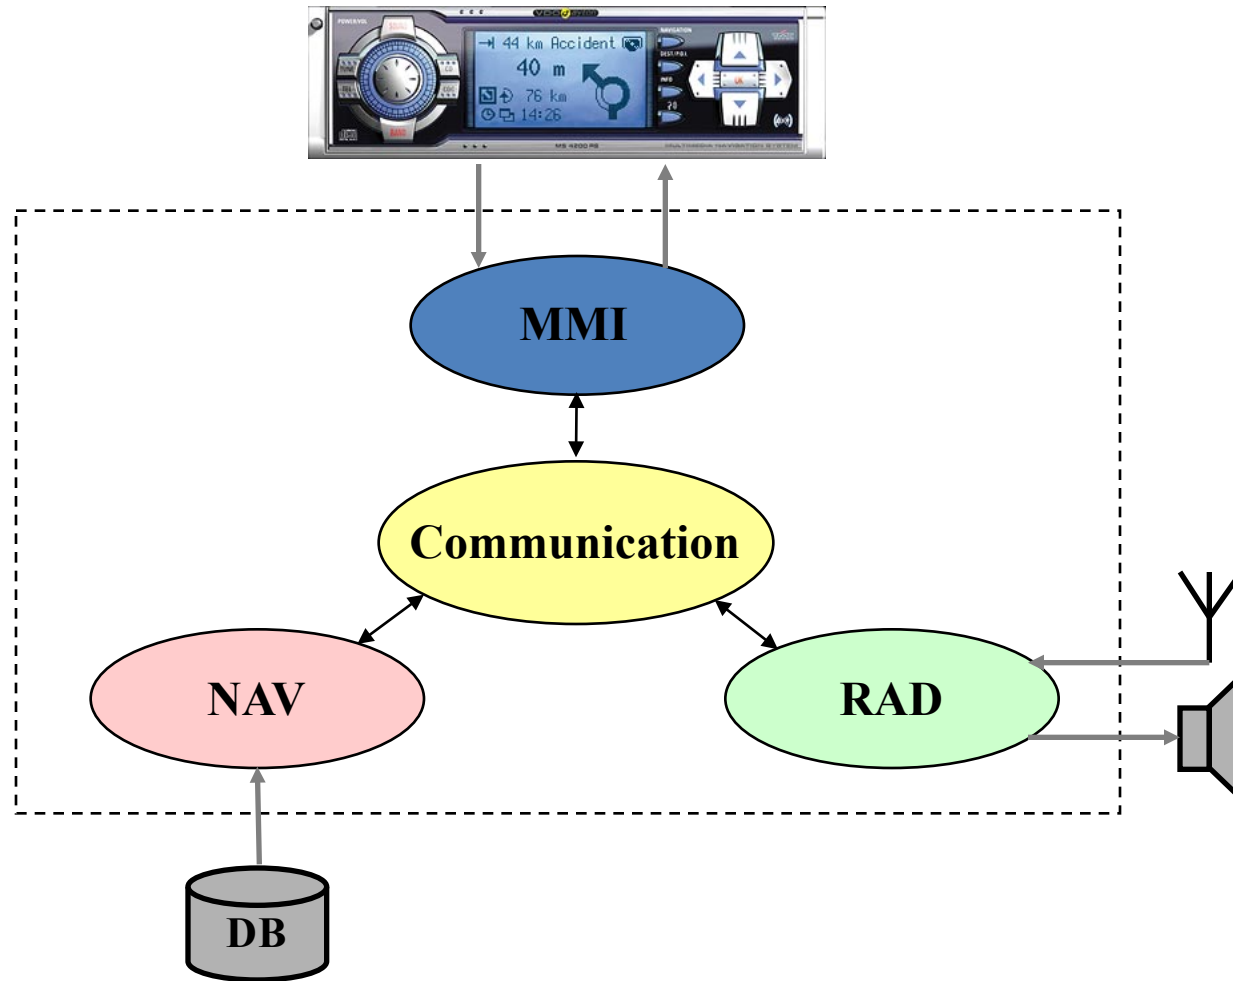
Example: In-Car Navigation System

- Car radio with navigation system
 - User interface should be responsive
 - Traffic messages (TMC) must be handled in a timely way
 - Several applications may execute concurrently



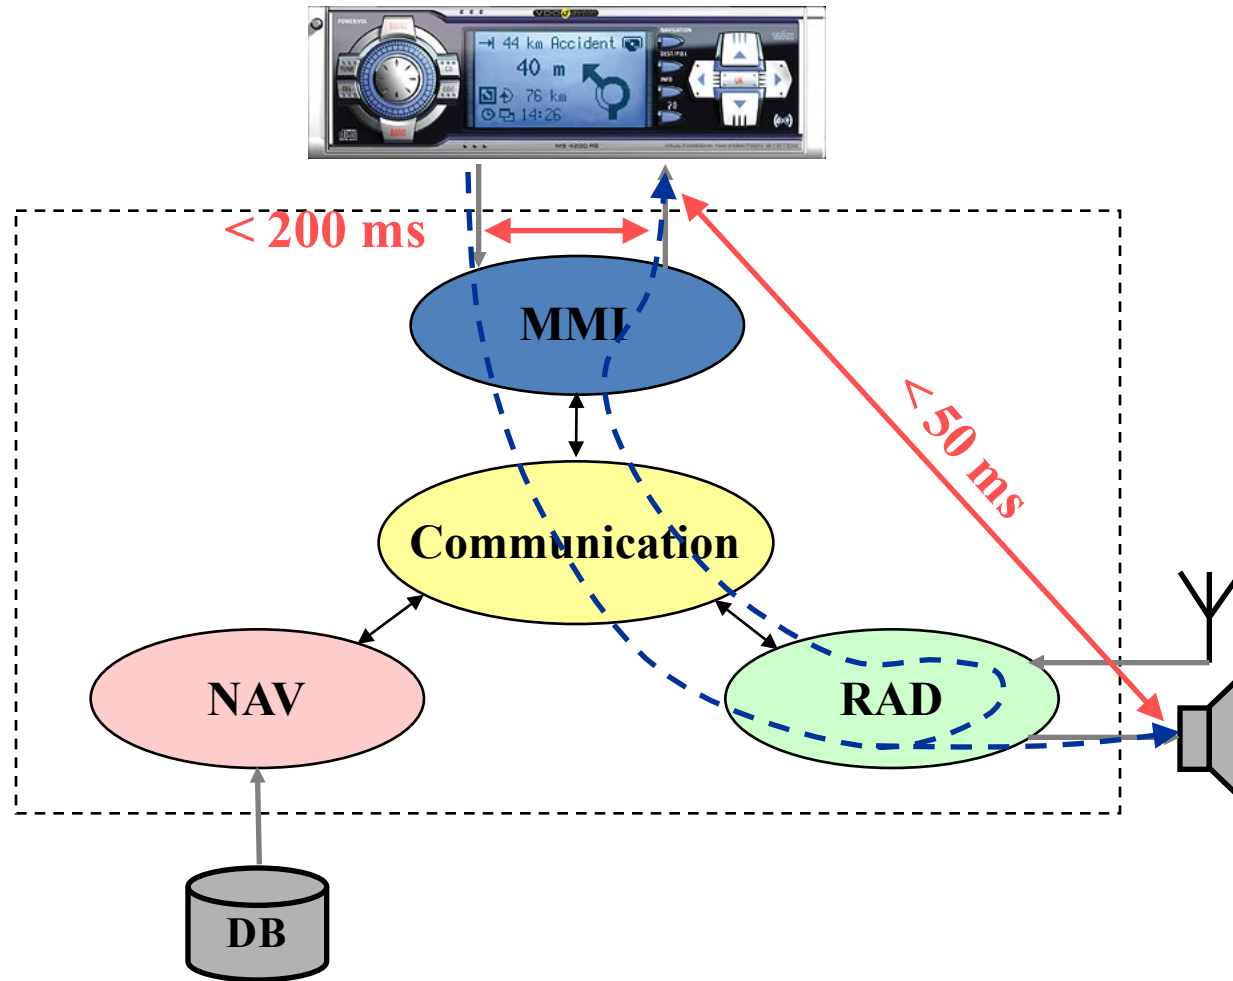
System Overview

© Thiele, ETHZ



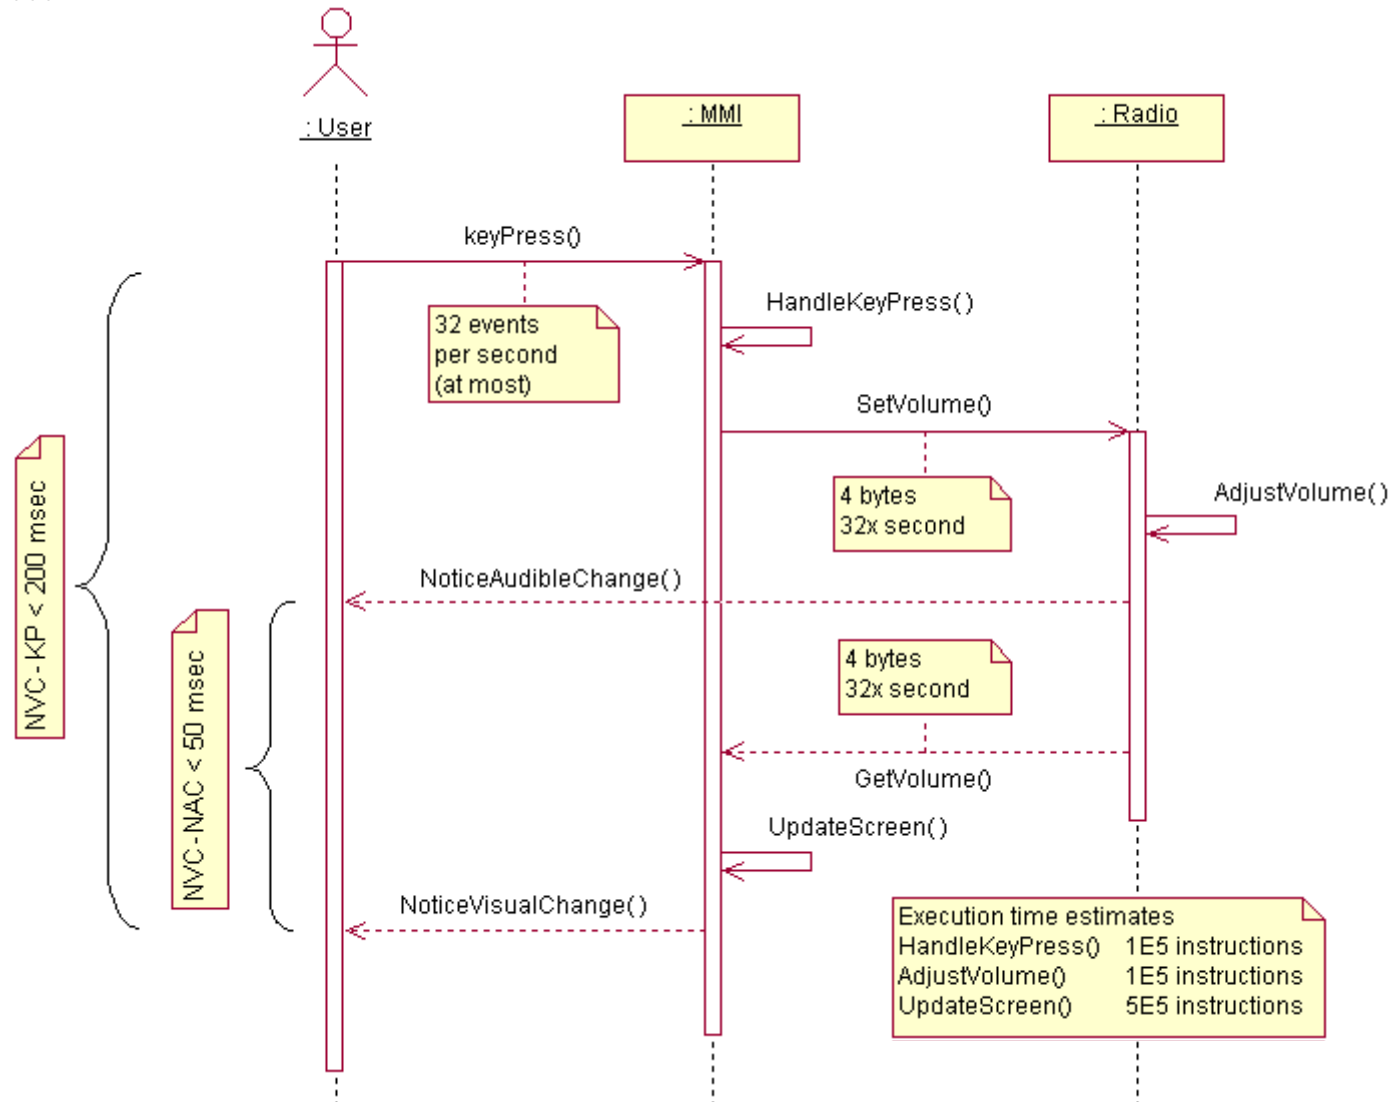
Use Case 1: Change Audio Volume

© Thiele, ETHZ



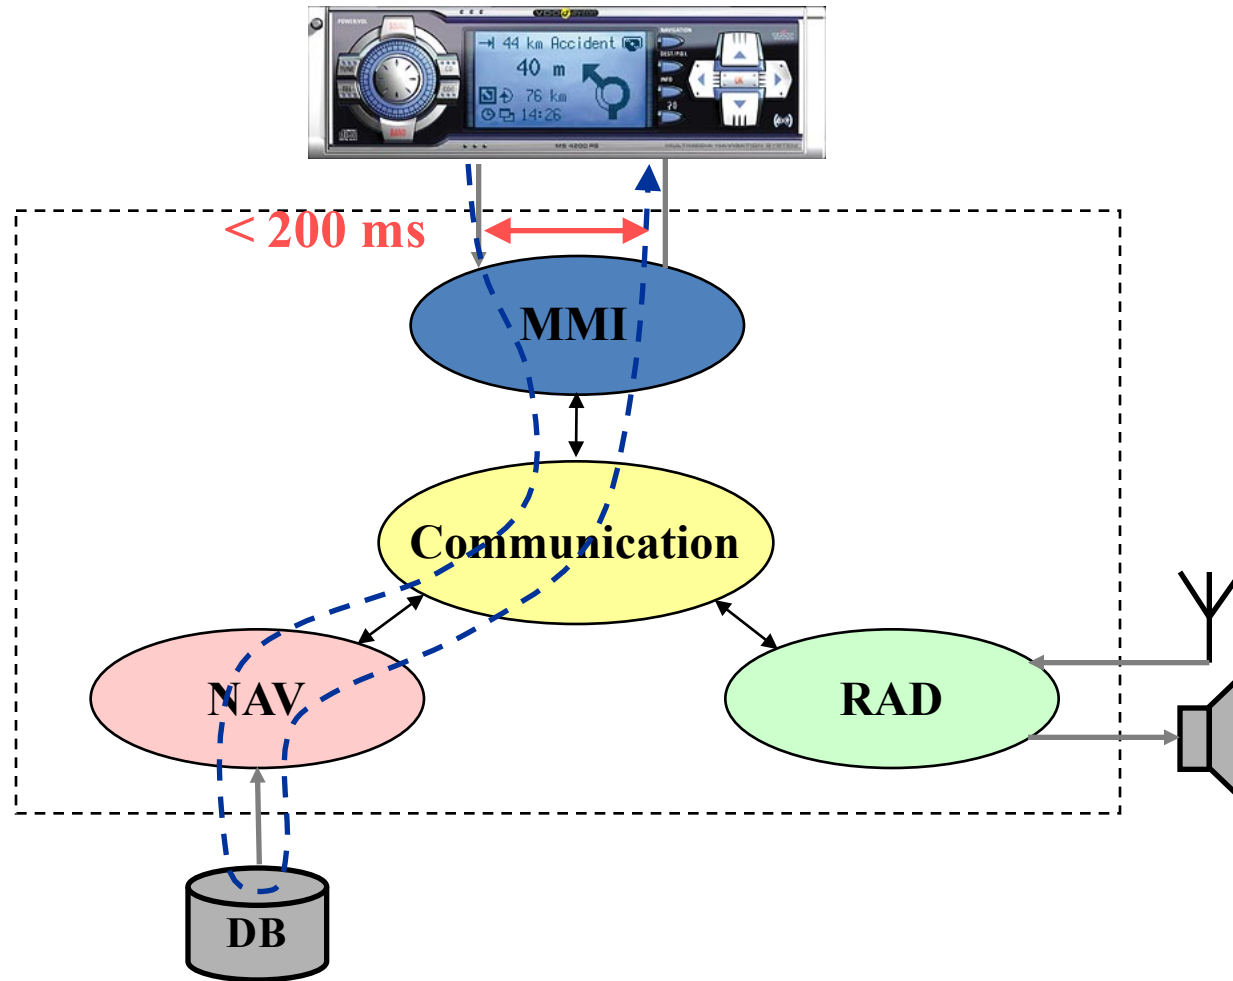
Use Case 1: Change Audio Volume

© Thiele, ETHZ




Use Case 2: Lookup Destination Address

© Thiele, ETHZ




Use Case 2: Lookup Destination Address

© Thiele, ETHZ



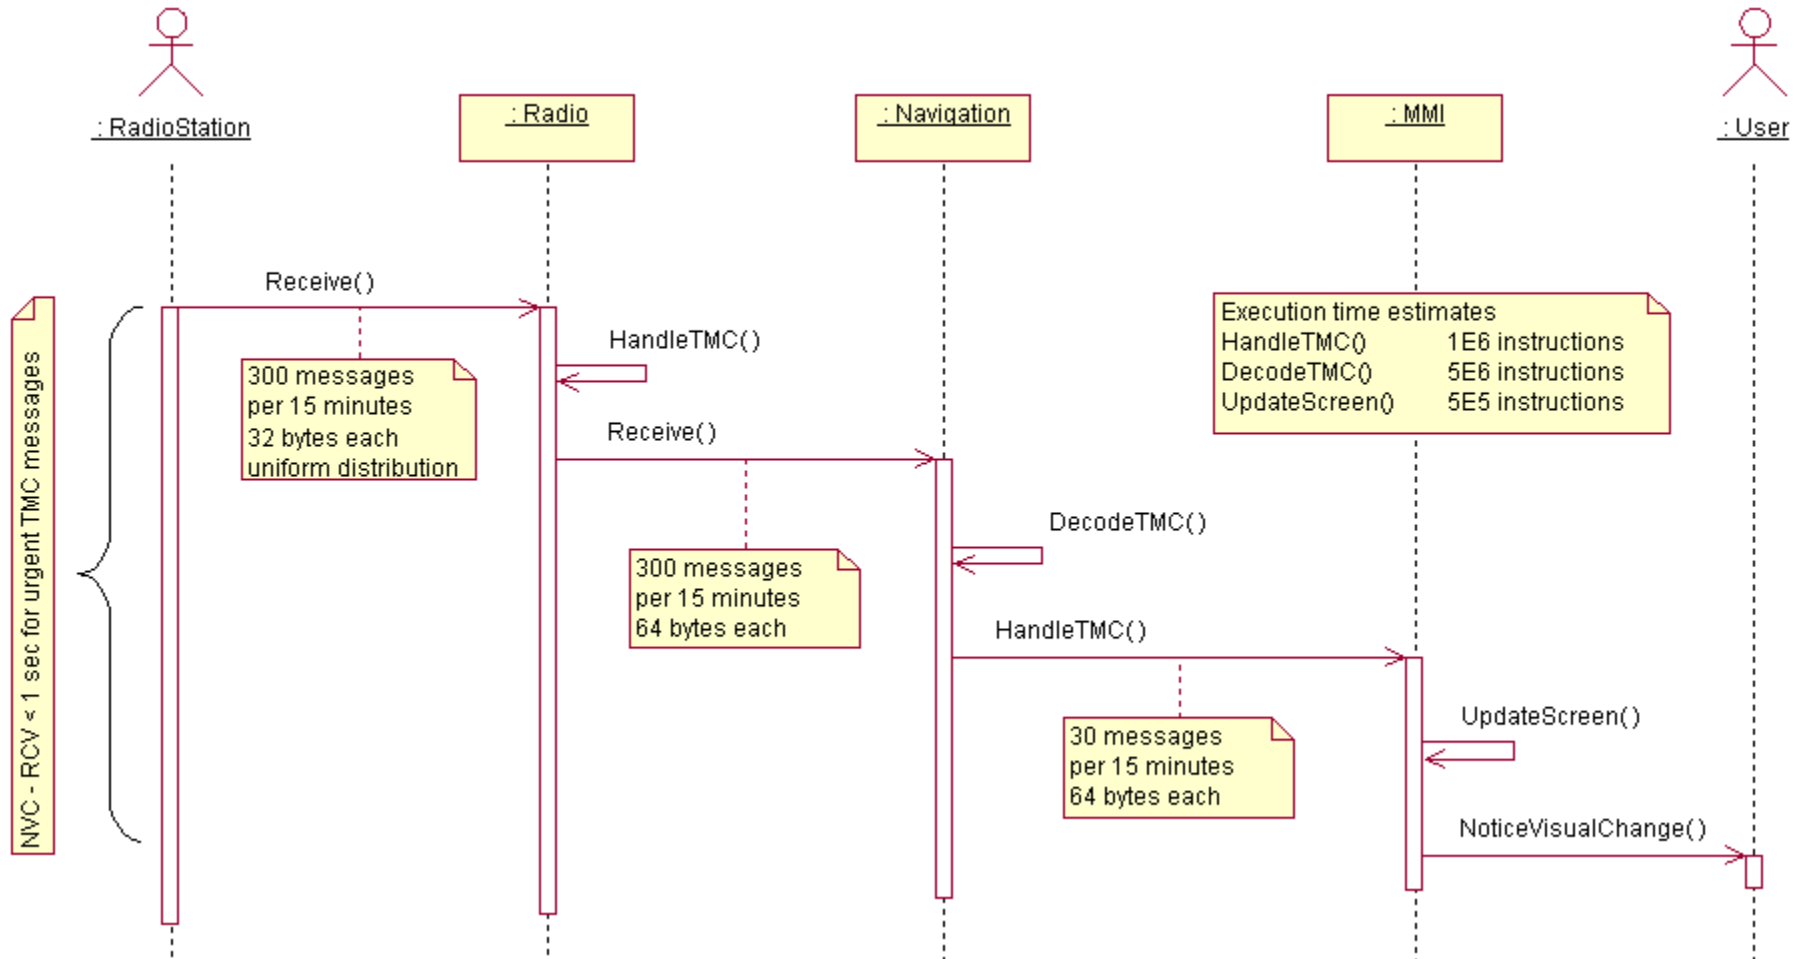
Use Case 3: Receive TMC Messages

© Thiele, ETHZ

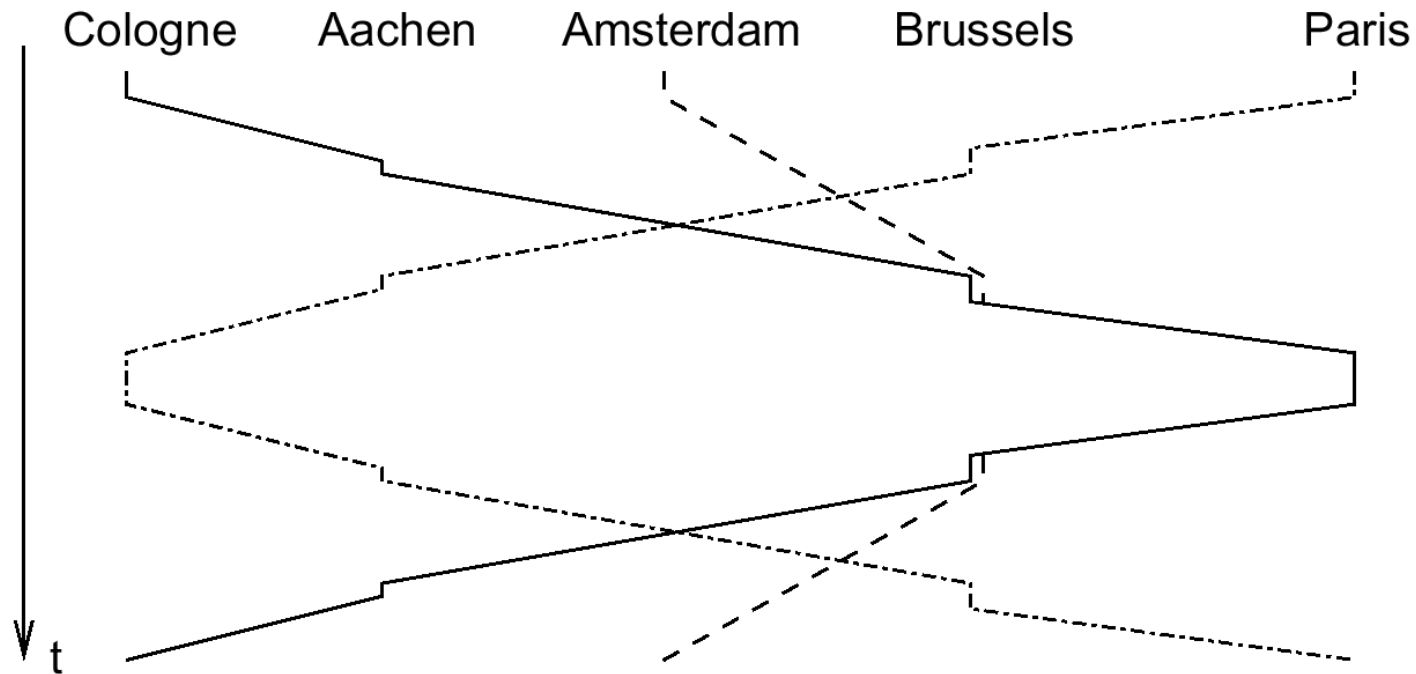


Use Case 3: Receive TMC Messages

© Thiele, ETHZ



Concurrency in MSC



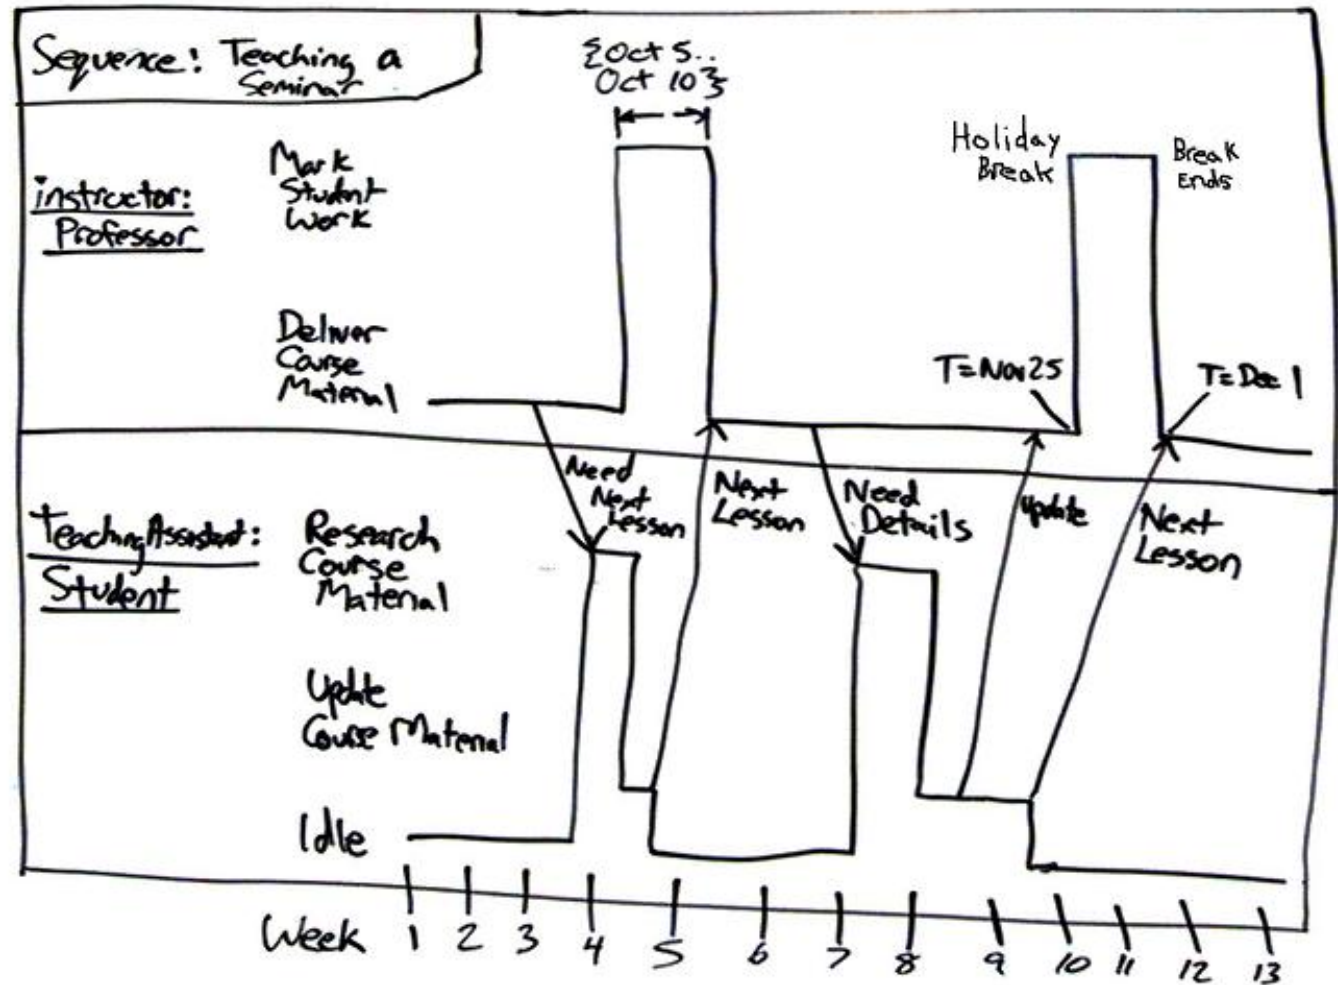
No distinction between accidental overlap and synchronization

- A special case of sequence charts



UML: Timing Diagrams

- Illustrate object state changes over time



© Scott Ambler,
Agile Modeling,
www.agilemodeling.com,
2003

(Message) Sequence Charts, Summarized

- **PROs**

- Appropriate for visualizing schedules
- Proven method for representing schedules in transportation
- Standard defined: *ITU-TS Recommendation Z.120: Message Sequence Chart (MSC)*, ITU-TS, Geneva, 1996
- Semantics also defined: *ITU-TS Recommendation Z.120: Message Sequence Chart (MSC)—Annex B: Algebraic Semantics of Message Sequence Charts*, ITU-TS, Geneva

- **CONS**

- Describes just one case!

Summary

- Introduction to requirements for specification languages
- Modeling Computation
 - Why not von Neumann?
 - Models for communication
 - Models of components
- Task graphs
- Supporting early design phases
 - Textual descriptions
 - Use cases or scenarios
 - (Message) sequence charts

Next Time

- Communicating Finite State Machines (CFSMs)
 - Chapter 2.4